


Homework 1


NetID: xm2074, jh7948
Name: Xiao Ma, Junru He

1. Determine if it is serializable of each executions.


a) . Serializable. The graph is asyclic.




b) . Not serializable. The graph has a cycle.



c) . Serializable. The graph is asyclic.



d) . Serializable. The graph is asyclic.



2. Proof of the serialization graph theorem.

Suppose the actual execution and the serial order have different reads-from. Suppose $R_m(x)$ of T_m reads x from $W_p(x)$ of T_p in the serial order, but $R_m(x)$ of T_m reads x from $W_q(x)$ of T_q in actual execution. Since both T_p and T_q contain writes to x , they must be connected in the serialization graph, so these three transactions must be ordered in the graph. Two cases:

1) $T_p \rightarrow T_q$. Then T_q must fall between T_p and T_m in serial schedule. So T_m reads-from T_q in serial schedule.

So $W_q(x)$ is the final write. This contradicts with the prerequisites that $W_p(x)$ is the final write in the serial order.

2) $T_q \rightarrow T_p$. Then T_m must read x from T_p in actual execution.

So $W_p(x)$ is the final write, which is contradicts the prerequisites that $W_q(x)$ is the final write in the actual order.

T1	T2
lock of x, write x	lock of y, write y
request lock of y -> waiting	request lock of x -> waiting

3. Is two phase locking deadlockfree?

No. It may produce a deadlock.

Considering the following scenario:

T1 has the lock of x and writes x, and T2 has the lock of y and writes y. Both of transaction T1 and T2 do not release the lock until the transaction ends. But at this time, T1 wants to read y and T2 wants to read x, so both of them request the lock of each other, and both of them owns the lock that each other needs. Then a deadlock produces.

4. “Pure” two phase locking -> order-preserving serializable executions.

To prove that “pure” two phase locking produces order-preserving serializable executions, we need two steps.

- First we need to prove that “pure” two phase locking produces serializable executions. That means there is no cycle in the serialization graph.

Suppose there is a cycle, take $T1 \rightarrow T2 \rightarrow T1$ as an example. For the definition, we know that locks can be released earlier than the end of the transaction but **no new locks are acquired** afterwards. That means the transaction finishes all possible conflict operations before releasing the lock. We show the time point that the transaction finishes its conflict operations and releases lock with $\text{lockpoint}(\text{transaction})$. For the example, $\text{lockpoint}(T1) < \text{lockpoint}(T2)$, that means T1 finishes all conflict operations and releases locks before T2 starts. Then from the graph, $\text{lockpoint}(T2) < \text{lockpoint}(T1)$, it contradicts with the statement before, so a cycle is impossible.

- Second we need to prove that “pure” two phase locking produces order-preserving serializable executions.

Take $T1 \rightarrow T2 \rightarrow T1$ as an example again. Suppose that T1 and T2 are not serializable. There must be some conflicts between T2 and T1 (conflict of T1 that happens after T2 begins). According to the definition of “pure” two phase locking, T2 begins when all conflicts in T1 have been finished. Then it is impossible that there is conflict of T1 that happens after T2 begins. So all operations of T1 must precede all operations of T2, and T1 and T2 is an equivalent serial execution.

From all above, “pure” two phase locking produces order-preserving serializable executions.

5. Would intention locking protocol be serializable?

No. The intention locking protocol is not serializable.

Considering the following scenario as the counterexample:

T1 requests intention write lock on Database and intention write lock on File 1. If intention lock can be released at any time, then T1 releases intention write lock on Database when T1 gets intention write lock on File 1.

At the same time, T2 requests read lock on Database and it succeeds.

Under this circumstance, T1 and T2 are not serializable.

6. Modify optimistic concurrency control algorithm.

We can record a priority according to the timestamp of every transaction. Then time stamp is the first time that the transaction begins and will never change even if it restarts.


The earlier the transaction starts, the priority is higher.

When there is a conflict, keep the transaction with higher priority and let it commit, then restart the other one. With this method, we can make sure that every transaction can eventually completes.


7. B-Tree

a)


Original data structure:



Insert 98, Split:



Delete 8, free-on-empty:



b)

```

begin at some non-leaf node n
while x is not in n:
  if x in keyset(n):
    return None # x is not found
  else if x is in inset(n) and edgeset(n, n'):
    n = n'
  else:
    set n to some ancestor node
if n is leaf:
  return key x and its value
else:
  return None
  
```

c)

For any search operation on any data item x beginning at any node m , the following invariant holds:

After any operation of any process, if the search/insert/delete for item x is at node n_1 , then either

item x is in $\text{keyset}(n_1)$ or there is a path from n_1 to node n_2 such that item x is in $\text{keyset}(n_2)$ and

every edge E along that path has x in its edgeset .

The invariant above works for ordinary B tree that data is stored not only in leaf nodes but also in

inner nodes. However, in our case, data is only stored in leave nodes and inner nodes only store indexes.

Therefore, in question (b), we make sure that the search algorithm gets valid result only on leaf nodes. So, the invariant also holds in the B-tree in our problem.

d)

When the node becomes a leaf node or all its children become empty.

8. Multi-version read consistency ensures serializability.

Considering the following scenario:

Suppose that the multi-version read consistency can not ensure the serializability, that means there is a cycle in the serializability graph. Then we suppose that $T_i \rightarrow T_j$ ($\text{lockpoint}(T_i) < \text{lockpoint}(T_j)$) has a conflict, it has three cases:

- Both T_i and T_j are read-only transaction:
It is serializable, because read can not conflicts.

- Ti is read-only and Tj is read-write:
In this case, Ti may read an item x before Tj write on x, and Tj writes on y before Ti reads item y. That means until Tj release its lock, Ti can read item y, so $\text{lockpoint}(Tj) < \text{lockpoint}(Ti)$, which contradicts with the assumption.
- Ti is read-write and Tj is read-only:
In this case, Ti holds a lock until Ti ends. After Ti releases the lock, Tj reads the item. Therefore, $\text{lockpoint}(Ti)$ precedes $\text{lockpoint}(Tj)$. It is serializable.

Above all, multi-version read consistency ensures serializability.

9. Oracle database system.

This can lead to a non-serializable execution.

Considering the following scenario:

T1: select * from db where id < 20

T2: insert into db values (10, 'Tom')

T1: select * from db where id < 25

For multi-version read consistency, T1 does not require locks. The first part of T1 executes before T2, and the second part executes after T2 releases locks. So the first part and the second part of T1 get different results from the same transaction. And the serializability graph is $T1 \rightarrow T2 \rightarrow T1$, having a cycle.

So it leads to a non-serializable execution.

10. AQuery

i) The code and the result are as following:

Code:

```
DROP TABLE IF EXISTS ticks

CREATE TABLE ticks(id INT, date INT, endofdayprice INT)

LOAD DATA INFILE "data/ticks.csv" INTO TABLE TICKS FIELDS
TERMINATED BY ","

SELECT id, max(ratios(endofdayprice)) as max_return,
min(ratios(endofdayprice)) as min_return FROM ticks
ASSUMING ASC id, ASC date
GROUP BY id
```

Result:

```
id | max_return | min_return
=====
3000 124 0.00909091
3001 117 0.00571429
3002 190 0.00869565
3003 83.5 0.037037
done.
```

ii)

Since id has been sorted, so group by id is an easy operation that cost little. Beside, another thing we need to do is to sort date for each id and calculate the ratio. Then we write the following code:

```
SELECT id, max(ratios(endofdayprice)) as max_return,
min(ratios(endofdayprice)) as min_return FROM ticks
ASSUMING ASC date
GROUP BY id
```